

Kaggle

Titanic: Machine Learning from Disaster

Introduction

This repository holds results for the Kaggle competition: Titanic: Machine Learning from Disaster.

Data

The datasets were obtained from the Kaggle Titanic Challenge: [Kaggle page](#)

- Training Dataset: [training data](#)
- Testing Dataset: [test data](#)

1. Loading Packages/ Data

```
for (package in c('knitr', 'caret', 'randomForest', 'e1071', 'gbm', 'rpart', 'rpart.plot', 'ggplot2', '')) {  
  if (!require(package, character.only = TRUE, quietly = FALSE)) {  
    install.packages(package)  
    library(package, character.only = TRUE)  
  }  
}  
  
val_dfname <- c("train.csv", "test.csv")  
val_dfpwd <- paste(getwd(), "/data", sep = "/")  
  
val_dtrawname <- c("data_training.raw", "data_testing.raw")  
val_dtname <- c("data_training", "data_testing")  
  
val_dtclass <- c("val_trainclass", "val_testclass")  
  
val_trainclass <- c("integer",    ## PassengerId  
                   "factor",     ## Survived  
                   "factor",     ## Pclass  
                   "character",  ## Name  
                   "factor",     ## Sex  
                   "numeric",    ## Age  
                   "integer",    ## SibSp  
                   "integer",    ## Parch  
                   "character",  ## Ticket  
                   "numeric",    ## Fare  
                   "character",  ## Cabin  
                   "factor")     ## Embarked  
  
val_testclass <- val_trainclass[-2]
```

```

for (i in 1:length(val_dtrawname)){

  assign(val_dtrawname[i], read.csv(paste(val_dfp, val_dfname[i], sep = "/"),
                                     na.strings = c("NA", ""),
                                     colClasses = get(val_dtclass[i])))

  assign(val_dtname[i], get(val_dtrawname[i]))

}

```

2. Pre-process the Data

Check the original data:

```

## dim(data_training.raw)
## str(data_training.raw)
summary(data_training.raw)

```

```

##   PassengerId   Survived  Pclass     Name       Sex
##   Min.   : 1.0   0:549    1:216  Length:891   female:314
##   1st Qu.:223.5  1:342    2:184   Class :character  male :577
##   Median :446.0           3:491   Mode  :character
##   Mean   :446.0
##   3rd Qu.:668.5
##   Max.   :891.0
##
##      Age          SibSp          Parch          Ticket
##   Min.   : 0.42   Min.   :0.000   Min.   :0.0000   Length:891
##   1st Qu.:20.12   1st Qu.:0.000   1st Qu.:0.0000   Class :character
##   Median :28.00   Median :0.000   Median :0.0000   Mode  :character
##   Mean   :29.70   Mean   :0.523   Mean   :0.3816
##   3rd Qu.:38.00   3rd Qu.:1.000   3rd Qu.:0.0000
##   Max.   :80.00   Max.   :8.000   Max.   :6.0000
##   NA's    :177
##      Fare          Cabin          Embarked
##   Min.   : 0.00   Length:891   C    :168
##   1st Qu.: 7.91   Class :character  Q    : 77
##   Median :14.45   Mode  :character  S    :644
##   Mean   :32.20           NA's: 2
##   3rd Qu.:31.00
##   Max.   :512.33
##

```

Categorize passengers by 'Title', and create new 'FamilySize' Variable:

```

for (i in 1:length(val_dtname)){

  temp_data <- get(val_dtname[i])
  temp_data["Title"] <- NA
  temp_data["FamilySize"] <- NA

  for (j in 1:nrow(temp_data)){

```

```

    temp_data[j, "Title"] <- strsplit(temp_data[j, "Name"], split=',.')[[1]][2]
    temp_data[j, "FamilySize"] <- temp_data[j, "SibSp"] + temp_data[j, "Parch"] + 1
  }

  temp_data[temp_data == ""] <- NA

  temp_data$Title = as.character(temp_data$Title)
  temp_data$FamilySize = as.integer(temp_data$FamilySize)
  ## print(sum(is.na(temp_data$Title)))
  ## print(sum(is.na(temp_data$FamilySize)))
  assign(val_dtname[i], temp_data)
}

rm(temp_data)

```

Replace NA values within numeric class columns with mean and NA values within other class columns with most common occurrence:

```

for (i in 1:length(val_dtname)){

  temp_data <- get(val_dtname[i])

  for (j in 1:ncol(temp_data)) {

    if (class(temp_data[, j]) == "numeric") {

      temp_colmean <- mean(temp_data[, j], na.rm = TRUE)
      temp_data[, j][which(is.na(temp_data[, j]))] <- temp_colmean

    } else {

      temp_colmode <- tail(names(sort(table(temp_data[, j]))), 1)
      temp_data[, j][which(is.na(temp_data[, j]))] <- temp_colmode

    }

  }

  assign(val_dtname[i], temp_data)
}

rm(temp_data, temp_colmean, temp_colmode)

```

Check the processed data:

```

## dim(data_training)
## str(data_training)
summary(data_training)

```

```
## PassengerId      Survived Pclass      Name      Sex
```

```
## Length:891      0:549      1:216      Length:891      female:314
## Class :character 1:342      2:184      Class :character male :577
## Mode :character      3:491      Mode :character
##
##
##
##      Age      SibSp      Parch      Ticket
## Min.   : 0.42      Length:891      Length:891      Length:891
## 1st Qu.:22.00      Class :character  Class :character  Class :character
## Median :29.70      Mode :character  Mode :character  Mode :character
## Mean   :29.70
## 3rd Qu.:35.00
## Max.   :80.00
##      Fare      Cabin      Embarked      Title
## Min.   : 0.00      Length:891      C:168      Length:891
## 1st Qu.: 7.91      Class :character  Q: 77      Class :character
## Median :14.45      Mode :character  S:646      Mode :character
## Mean   :32.20
## 3rd Qu.:31.00
## Max.   :512.33
##      FamilySize
## Length:891
## Class :character
## Mode :character
##
##
##
```

Check the processed data:

```
tblsumfunc <- function(x){

  temp_data <- data.frame(Survived = data_training$Survived, Title = data_training[[x]], stringsAsFactors = FALSE)
  temp_obscount <- sort(table(temp_data[, 2]), decreasing = FALSE)

  if (nrow(temp_obscount) > 10) {

    if (class(temp_data[, 2]) == "numeric") {

      temp_data[, 2] <- 10 * ceiling(temp_data[, 2] / 10)
      ## table(temp_data)

    } else {

      temp_lfobsnm <- names(temp_obscount[1:(dim(temp_obscount) - 10)])
      temp_data[, 2][which(is.element(temp_data[, 2], temp_lfobsnm))] <- "Other"
      ## table(temp_data)

    }

  }

  temp_table <- table(temp_data)
  temp_sumtable <- addmargins(temp_table, FUN = list(Total = sum), quiet = TRUE)
```

```

temp_proptable <- prop.table(temp_sumtable[c(1, 2),], 2)
temp_mergedtable <- rbind(temp_sumtable[1, ],
                           temp_proptable[1, ],
                           temp_sumtable[2, ],
                           temp_proptable[2, ],
                           temp_sumtable[3, ])
rownames(temp_mergedtable) <- c("Didn't Survive", "%", "Survived", "%", "Total")

print(x)
temp_kabletable <- kable(temp_mergedtable, digits = 2, caption = "test", output = FALSE)
cat(temp_kabletable, sep="\n")
cat(sep="\n\n")

rm(temp_data, temp_table, temp_sumtable, temp_proptable, temp_mergedtable)
}

val_sumcolname <- list("Pclass", "Title", "Sex", "Age", "FamilySize")

for(colname in val_sumcolname) { tblsumfunc(colname) }

```

```
## [1] "Pclass"
```

```
## Table: test
```

```
##
##           1           2           3      Total
## -----
## Didn't Survive    80.00    97.00    372.00    549.00
## %                 0.37     0.53     0.76     0.62
## Survived         136.00    87.00    119.00    342.00
## %                 0.63     0.47     0.24     0.38
## Total            216.00    184.00    491.00    891.00
##
```

```
## [1] "Title"
```

```
## Table: test
```

```
##
##           Col      Dr      Major      Master      Miss      Mlle      Mr      Mrs      Rev      the Cou
## -----
## Didn't Survive    1.0    4.00      1.0    17.00    55.0         0   436.00    26.00      6
## %                 0.5    0.57      0.5     0.42     0.3         0     0.84     0.21      1
## Survived          1.0    3.00      1.0    23.00   127.0         2    81.00    99.00      0
## %                 0.5    0.43      0.5     0.57     0.7         1     0.16     0.79      0
## Total             2.0    7.00      2.0    40.00   182.0         2   517.00   125.00      6
##
```

```
## [1] "Sex"
```

```
## Table: test
```

```
##
##           female      male      Total
## -----
## Didn't Survive    81.00    468.00    549.00
## %                 0.26     0.81     0.62
## Survived         233.00    109.00    342.00
## %                 0.74     0.19     0.38
## Total            314.00    577.00    891.00

```

```
##
## [1] "Age"
## Table: test
##
##      10      20      30      40      50      60      70      80      Total
## -----
## Didn't Survive 26.00  71.00  271.00  86.00  53.00  25.0  13.00  4.0  549.00
## %              0.41   0.62   0.67   0.55   0.62   0.6   0.76  0.8   0.62
## Survived      38.00  44.00  136.00  69.00  33.00  17.0   4.00  1.0  342.00
## %              0.59   0.38   0.33   0.45   0.38   0.4   0.24  0.2   0.38
## Total         64.00  115.00  407.00  155.00  86.00  42.0  17.00  5.0  891.00
##
## [1] "FamilySize"
## Table: test
##
##      1  11      2      3      4      5      6      7      8      Total
## -----
## Didn't Survive 374.0  7  72.00  43.00  8.00  12.0  19.00  8.00  6  549.00
## %              0.7   1  0.45  0.42  0.28  0.8   0.86  0.67  1  0.62
## Survived      163.0  0  89.00  59.00  21.00  3.0   3.00  4.00  0  342.00
## %              0.3   0  0.55  0.58  0.72  0.2   0.14  0.33  0  0.38
## Total         537.0  7  161.00  102.00  29.00  15.0  22.00  12.00  6  891.00
```



At a high level, the data suggests that passengers within the following groups had an improved survival rate:
 * Were Class 1 passengers * Had a title of 'Master' / aged 0-10 * Were female * Boarded with a family of size 2-4

Chart the processed data:

```
val_agehist <- ggplot(data_training, aes(x = Age, fill = Survived)) +
  geom_histogram() +
  ggtitle("Age vs Survival") +
  theme(legend.position = "bottom") +
  scale_fill_discrete(labels = c("No", "Yes"))

val_sexhist <- ggplot(data_training, aes(x = Sex, fill = Survived)) +
  geom_histogram() +
  ggtitle("Age vs Survival") +
  theme(legend.position = "bottom") +
  scale_fill_discrete(labels = c("No", "Yes"))

grid.arrange(val_agehist, val_sexhist, ncol = 2)
```

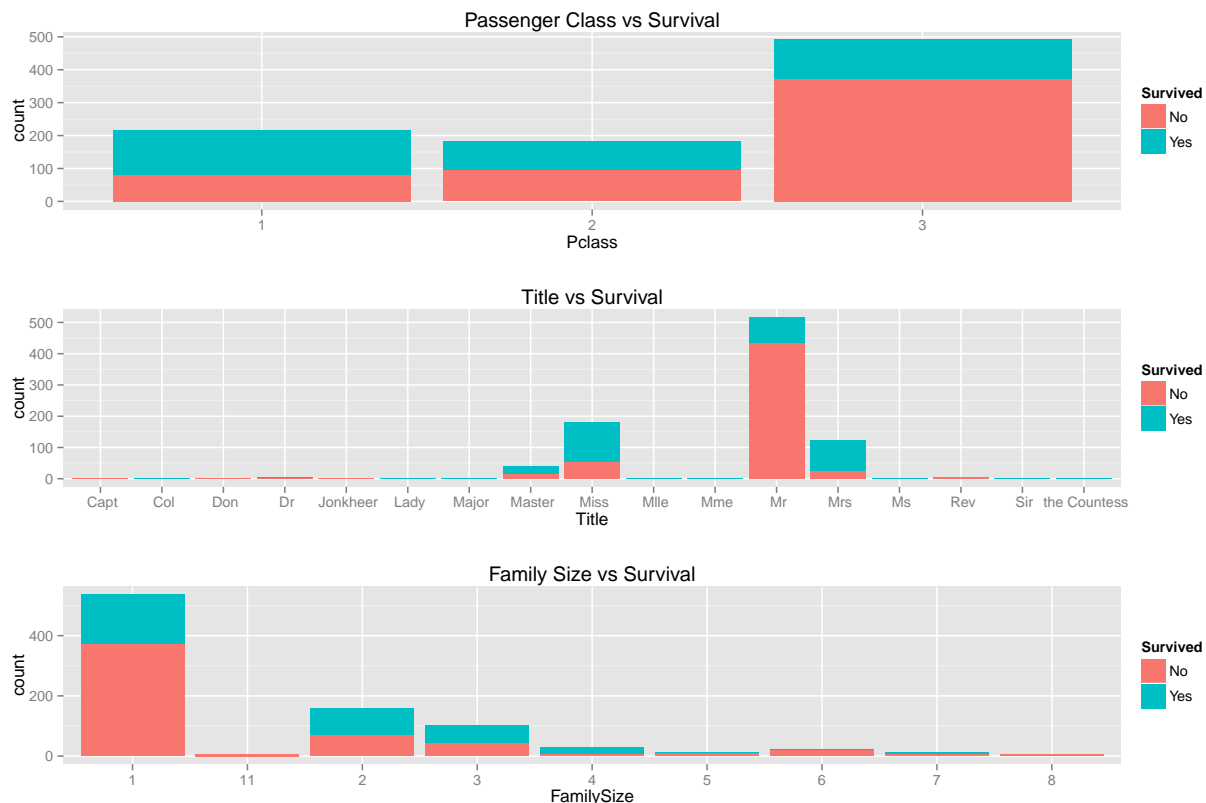


```
val_pclasshist <- ggplot(data_training, aes(x = Pclass, fill = Survived)) +
  geom_histogram() +
  ggtitle("Passenger Class vs Survival") +
  theme(legend.position = "right") +
  scale_fill_discrete(labels = c("No", "Yes"))

val_titlehist <- ggplot(data_training, aes(x = Title, fill = Survived)) +
  geom_histogram() +
  ggtitle("Title vs Survival") +
  theme(legend.position = "right") +
  scale_fill_discrete(labels = c("No", "Yes"))

val_familyhist <- ggplot(data_training, aes(x = FamilySize, fill = Survived)) +
  geom_histogram() +
  ggtitle("Family Size vs Survival") +
  theme(legend.position = "right") +
  scale_fill_discrete(labels = c("No", "Yes"))

grid.arrange(val_pclasshist, val_titlehist, val_familyhist, nrow = 3)
```



3. Prediction Modelling

Split the training data:

```
set.seed(12345)
data_training.rows <- createDataPartition(data_training$Survived, p = 0.7, list = FALSE)

data_training.train <- data_training[data_training.rows, ]
data_training.test <- data_training[-data_training.rows, ]
```

Check the split data:

```
## dim(data_training.train)
## str(data_training.train)

## 'data.frame':  625 obs. of  14 variables:
## $ PassengerId: chr  "1" "3" "4" "6" ...
## $ Survived   : Factor w/ 2 levels "0","1": 1 2 2 1 1 1 2 2 2 1 ...
## $ Pclass     : Factor w/ 3 levels "1","2","3": 3 3 1 3 1 3 3 2 1 3 ...
## $ Name       : chr  "Braund, Mr. Owen Harris" "Heikkinen, Miss. Laina" "Futrelle, Mrs. Jacques Heat
## $ Sex        : Factor w/ 2 levels "female","male": 2 1 1 2 2 2 1 1 1 2 ...
## $ Age        : num  22 26 35 29.7 54 ...
## $ SibSp      : chr  "1" "0" "1" "0" ...
## $ Parch      : chr  "0" "0" "0" "0" ...
## $ Ticket     : chr  "A/5 21171" "STON/O2. 3101282" "113803" "330877" ...
```



```
## $ Fare      : num  7.25 7.92 53.1 8.46 51.86 ...
## $ Cabin     : chr   "G6" "G6" "C123" "G6" ...
## $ Embarked  : Factor w/ 3 levels "C","Q","S": 3 3 3 2 3 3 3 1 3 3 ...
## $ Title     : chr   " Mr" " Miss" " Mrs" " Mr" ...
## $ FamilySize : chr   "2" "1" "2" "1" ...
```

```
## summary(data_training.train)
```

```
## dim(data_training.test)
str(data_training.test)
```

```
## 'data.frame': 266 obs. of 14 variables:
## $ PassengerId: chr  "2" "5" "11" "16" ...
## $ Survived   : Factor w/ 2 levels "0","1": 2 1 2 2 2 1 1 2 1 2 ...
## $ Pclass     : Factor w/ 3 levels "1","2","3": 1 3 3 2 3 2 3 3 2 3 ...
## $ Name       : chr  "Cumings, Mrs. John Bradley (Florence Briggs Thayer)" "Allen, Mr. William Henry"
## $ Sex        : Factor w/ 2 levels "female","male": 1 2 1 1 1 2 1 1 2 2 ...
## $ Age        : num   38 35 4 55 29.7 ...
## $ SibSp      : chr   "1" "0" "1" "0" ...
## $ Parch      : chr   "0" "0" "1" "0" ...
## $ Ticket     : chr  "PC 17599" "373450" "PP 9549" "248706" ...
## $ Fare       : num   71.28 8.05 16.7 16 7.22 ...
## $ Cabin      : chr  "C85" "G6" "G6" "G6" ...
## $ Embarked   : Factor w/ 3 levels "C","Q","S": 1 3 3 3 1 3 3 3 3 1 ...
## $ Title      : chr   " Mrs" " Mr" " Miss" " Mrs" ...
## $ FamilySize : chr  "2" "1" "3" "1" ...
```

```
## summary(data_training.test)
```

```
set.seed(12345)
val_dtmodel <- rpart(Survived ~ Pclass + Sex + Age + Fare + Embarked + FamilySize, data = data_training)
val_dtmodel.predict <- predict(val_dtmodel, data_training.test, type = "class")
val_dtcm <- confusionMatrix(val_dtmodel.predict, data_training.test$Survived)
val_dtcm
```

Decision tree prediction

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 145  38
##           1  19  64
##
##           Accuracy : 0.7857
##           95% CI : (0.7315, 0.8335)
##           No Information Rate : 0.6165
##           P-Value [Acc > NIR] : 2.603e-09
##
##           Kappa : 0.5303
```

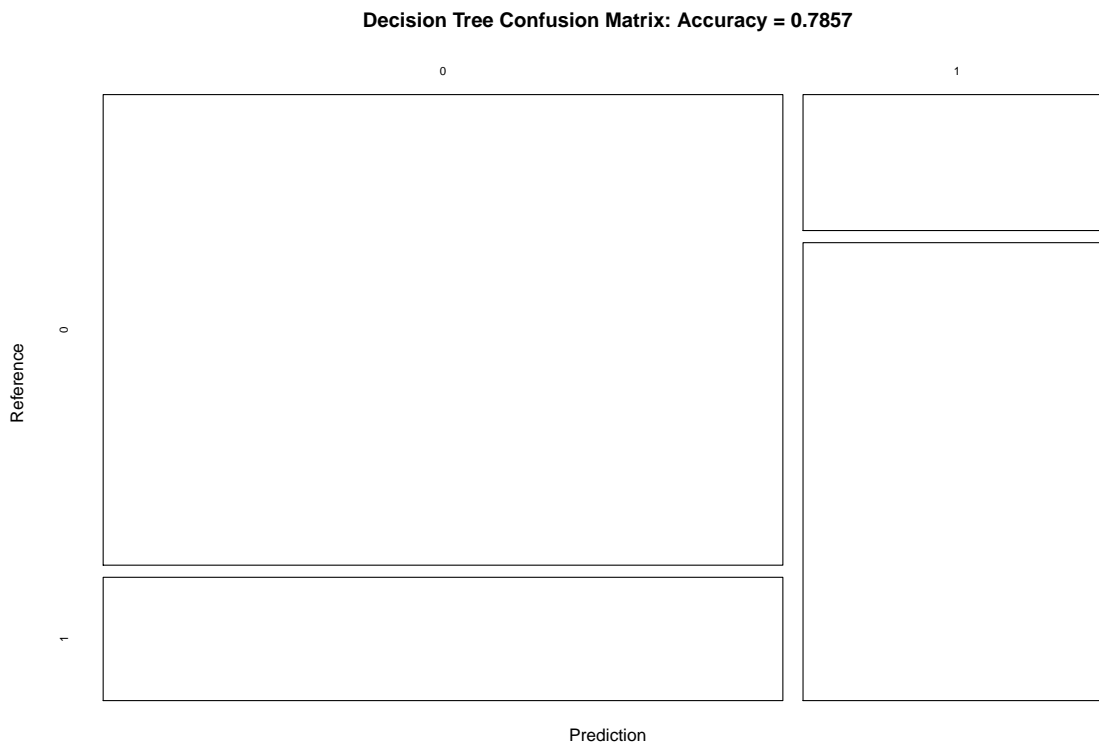
```
## McNemar's Test P-Value : 0.01712
##
##      Sensitivity : 0.8841
##      Specificity : 0.6275
##      Pos Pred Value : 0.7923
##      Neg Pred Value : 0.7711
##      Prevalence : 0.6165
##      Detection Rate : 0.5451
##      Detection Prevalence : 0.6880
##      Balanced Accuracy : 0.7558
##
##      'Positive' Class : 0
##
```

Decision tree prediction has a reported accuracy against the training dataset:

```
round(val_dtc$overall['Accuracy'], 4)
```

```
## Accuracy
## 0.7857
```

```
plot(val_dtc$table,
     col = val_dtc$byClass,
     main = paste("Decision Tree Confusion Matrix: Accuracy =",
     round(val_dtc$overall['Accuracy'], 4)))
```



```

set.seed(12345)
val_rfmodel <- randomForest(Survived ~ Pclass + Sex + Age + Fare + Embarked + FamilySize, data = data_t
val_rfmodel.predict <- predict(val_rfmodel, data_training.test, type = "class")
val_rfcM <- confusionMatrix(val_rfmodel.predict, data_training.test$Survived)
val_rfcM

```

Random forest prediction

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 149  33
##              1  15  69
##
##              Accuracy : 0.8195
##              95% CI : (0.768, 0.8638)
##      No Information Rate : 0.6165
##      P-Value [Acc > NIR] : 5.675e-13
##
##              Kappa : 0.6052
##  McNemar's Test P-Value : 0.01414
##
##              Sensitivity : 0.9085
##              Specificity : 0.6765
##              Pos Pred Value : 0.8187
##              Neg Pred Value : 0.8214
##              Prevalence : 0.6165
##              Detection Rate : 0.5602
##      Detection Prevalence : 0.6842
##              Balanced Accuracy : 0.7925
##
##              'Positive' Class : 0
##

```

Random forest prediction has a reported accuracy against the training dataset:

```

round(val_rfcM$overall['Accuracy'], 4)

```

```


## Accuracy
##    0.8195

```

```

plot(val_rfcM$table,
     col = val_rfcM$byClass,
     main = paste("Random Forest Confusion Matrix: Accuracy =",
     round(val_rfcM$overall['Accuracy'], 4)))

```



```
set.seed(12345)
val_fitControl <- trainControl(method = "repeatedcv", number = 5, repeats = 1)
val_gbmmodel <- train(Survived ~ Pclass + Sex + Age + Fare + Embarked + FamilySize, data = data_training,
  method = "gbm", control = val_fitControl)
val_gbmmodel.predict <- predict(val_gbmmodel, newdata = data_training.test)
val_gbmcm <- confusionMatrix(val_gbmmodel.predict, data_training.test$Survived)
val_gbmcm
```

Generalized boosted regression prediction

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 148  39
##           1  16  63
##
##           Accuracy : 0.7932
##           95% CI : (0.7395, 0.8403)
##           No Information Rate : 0.6165
##           P-Value [Acc > NIR] : 4.694e-10
##
##           Kappa : 0.5432
##           McNemar's Test P-Value : 0.003012
##
```

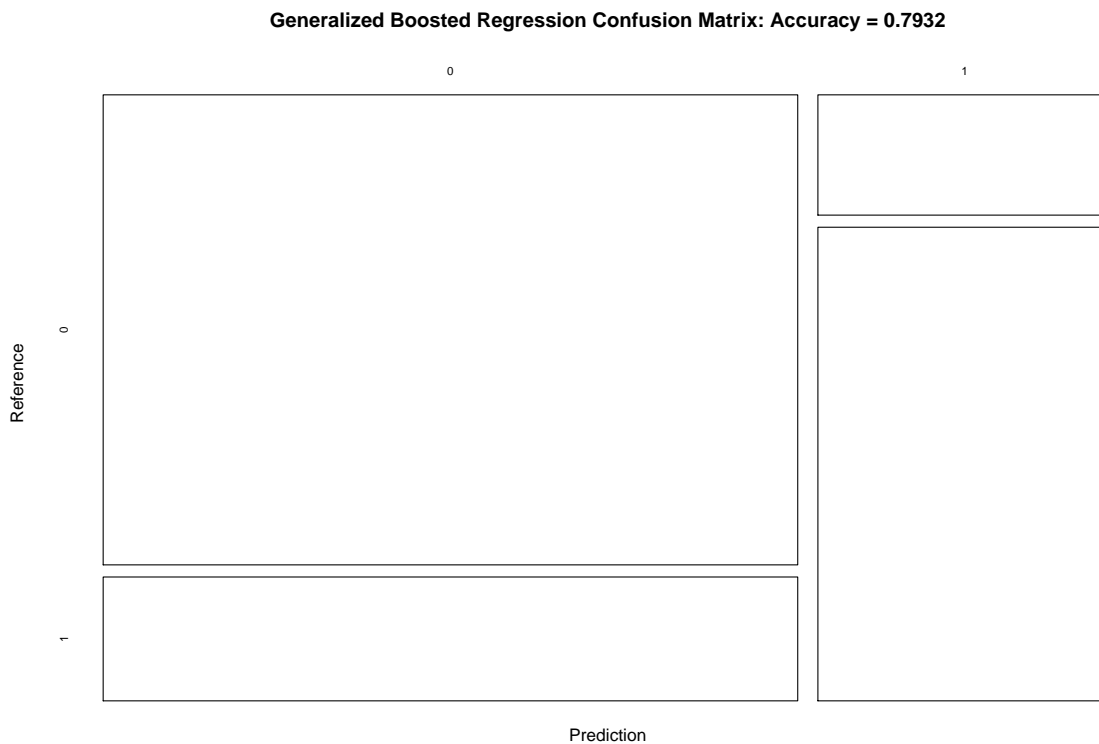
```
##          Sensitivity : 0.9024
##          Specificity : 0.6176
##          Pos Pred Value : 0.7914
##          Neg Pred Value : 0.7975
##          Prevalence : 0.6165
##          Detection Rate : 0.5564
##          Detection Prevalence : 0.7030
##          Balanced Accuracy : 0.7600
##
##          'Positive' Class : 0
##
```

Generalized boosted regression prediction has a reported accuracy against the training dataset:

```
round(val_gbmcm$overall['Accuracy'], 4)
```

```
## Accuracy
##    0.7932
```

```
plot(val_gbmcm$table,
     col = val_gbmcm$byClass,
     main = paste("Generalized Boosted Regression Confusion Matrix: Accuracy =",
                  round(val_gbmcm$overall['Accuracy'], 4)))
```



4. Model Selection

The expected out-of-sample error is calculated as 1 - accuracy for predictions made against the cross-validation set:

```
val_ooserror <- 1 - round(val_rfcm$overall['Accuracy'], 4)
## val_ooserror <- 1 - round(val_gbmcm$overall['Accuracy'], 4)
val_ooserror
```

```
## Accuracy
##    0.1805
```

```
val_selmodel.final <- predict(val_rfmodel, data_testing)
## val_selmodel.final <- predict(val_gbmmodel, data_testing)
```

5. Kaggle Submission

```
data_prediction <- data.frame(PassengerId = data_testing$PassengerId, Survived = val_selmodel.final)
write.table(data_prediction, "data/prediction.csv", row.names = FALSE, sep=",", col.names = TRUE)
```